

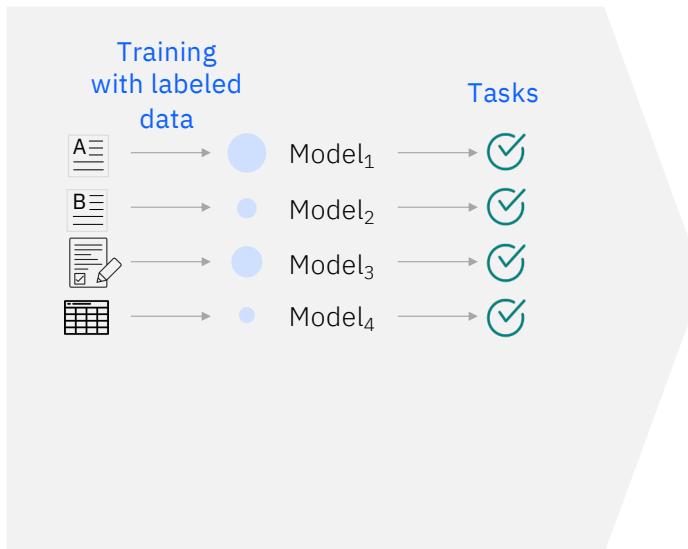
# Introduction to key concepts

# Agenda

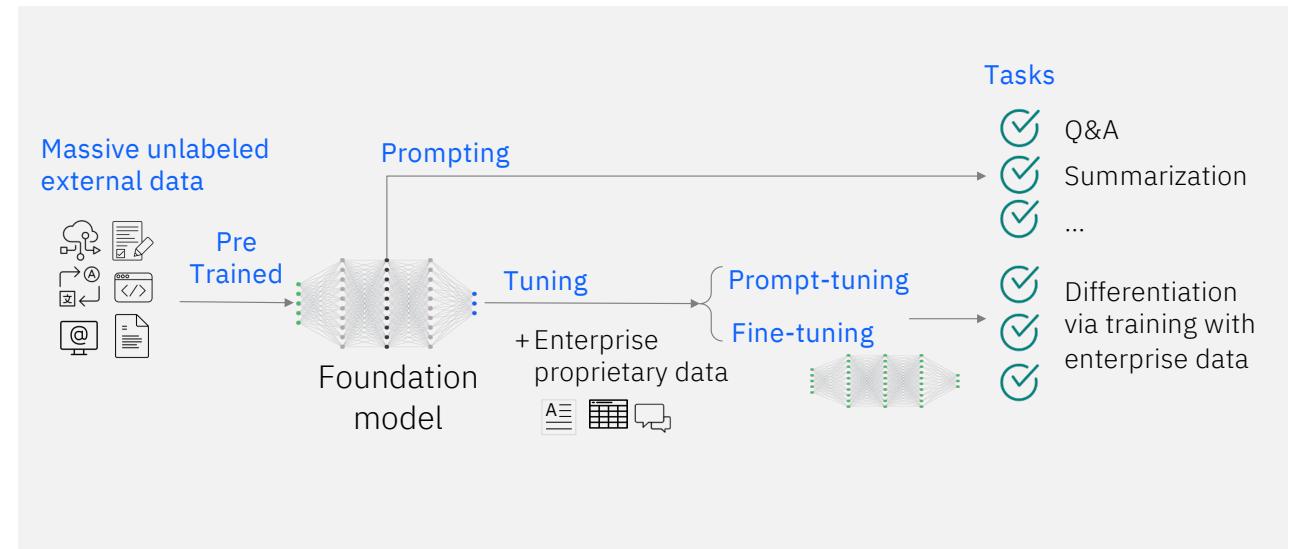
- FM, Prompt Engineering, Prompt Tuning and Fine Tuning
- Use Gen-AI API from python code in Notebooks or VS Code
- Leverage LangChain
- Retrieval Augmented Generation
- Introduction to Embeddings and VectorDB
- Position IBM Watson & other products as relevant

# Foundational models enable a new paradigm of data-efficient AI development – generative AI

## Traditional AI models



## Foundation Models

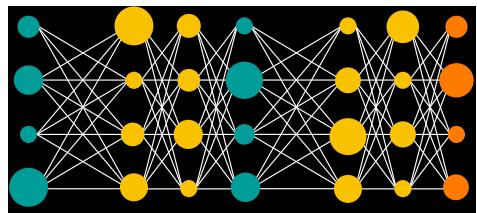


- Individual siloed models
- Require task specific training
- Lots of human supervised training

- Rapid adaptation to multiple tasks with small amounts of task-specific data
- Pre-trained unsupervised learning

# Foundation models are ...

 Self-supervised training



Foundation model

## Pre-trained

On unlabeled datasets of different modalities (e.g., language, time-series, tabular)

## Self-learning

Systems that leverage self-supervised learning

## Multiple applications

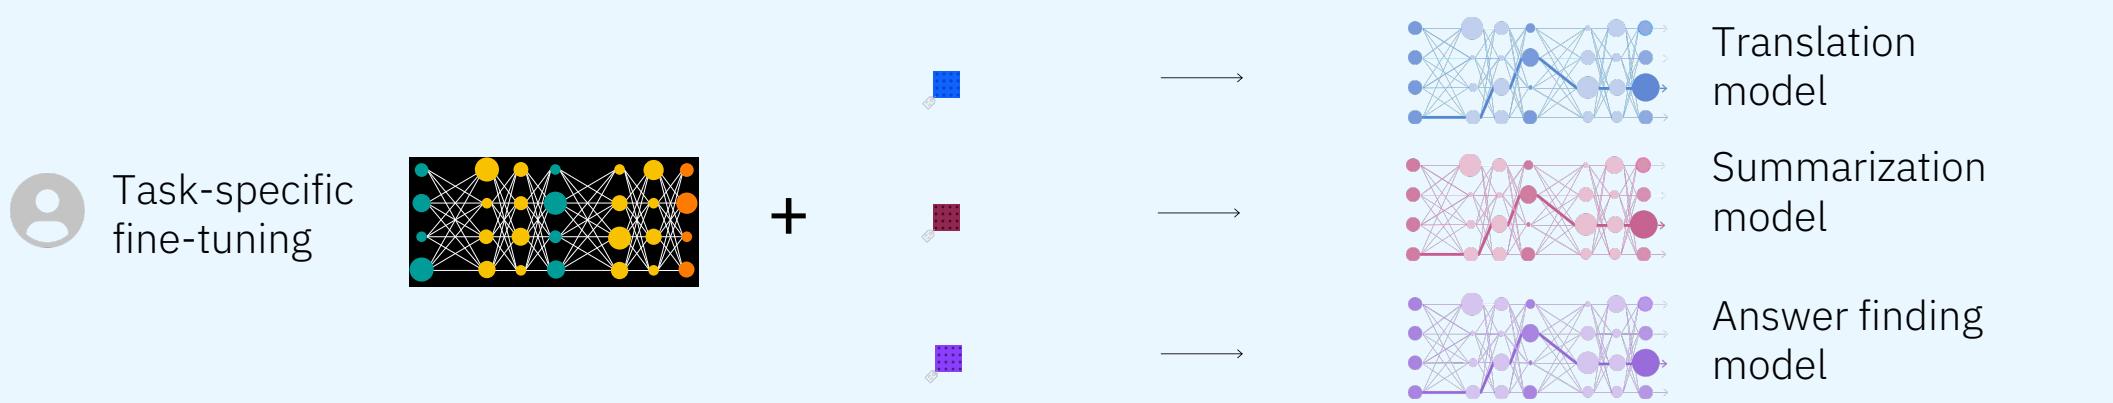
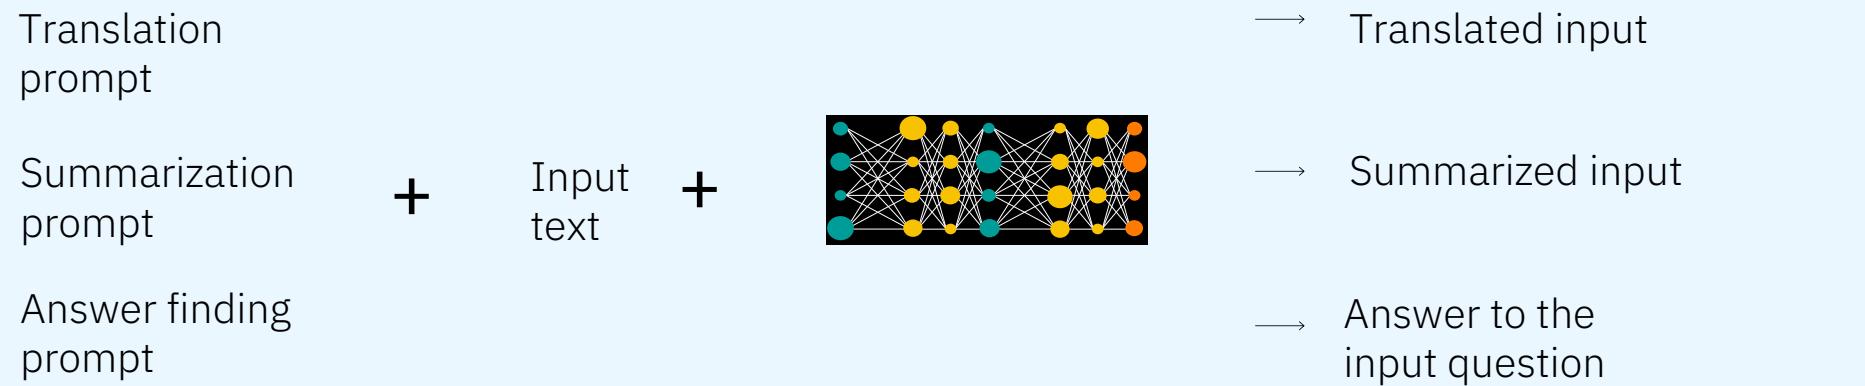
Able to learn generalizable and adaptable data representations that can be effectively used in a variety of domains and tasks (code generation, question answering, sentiment analysis)

## Large language models

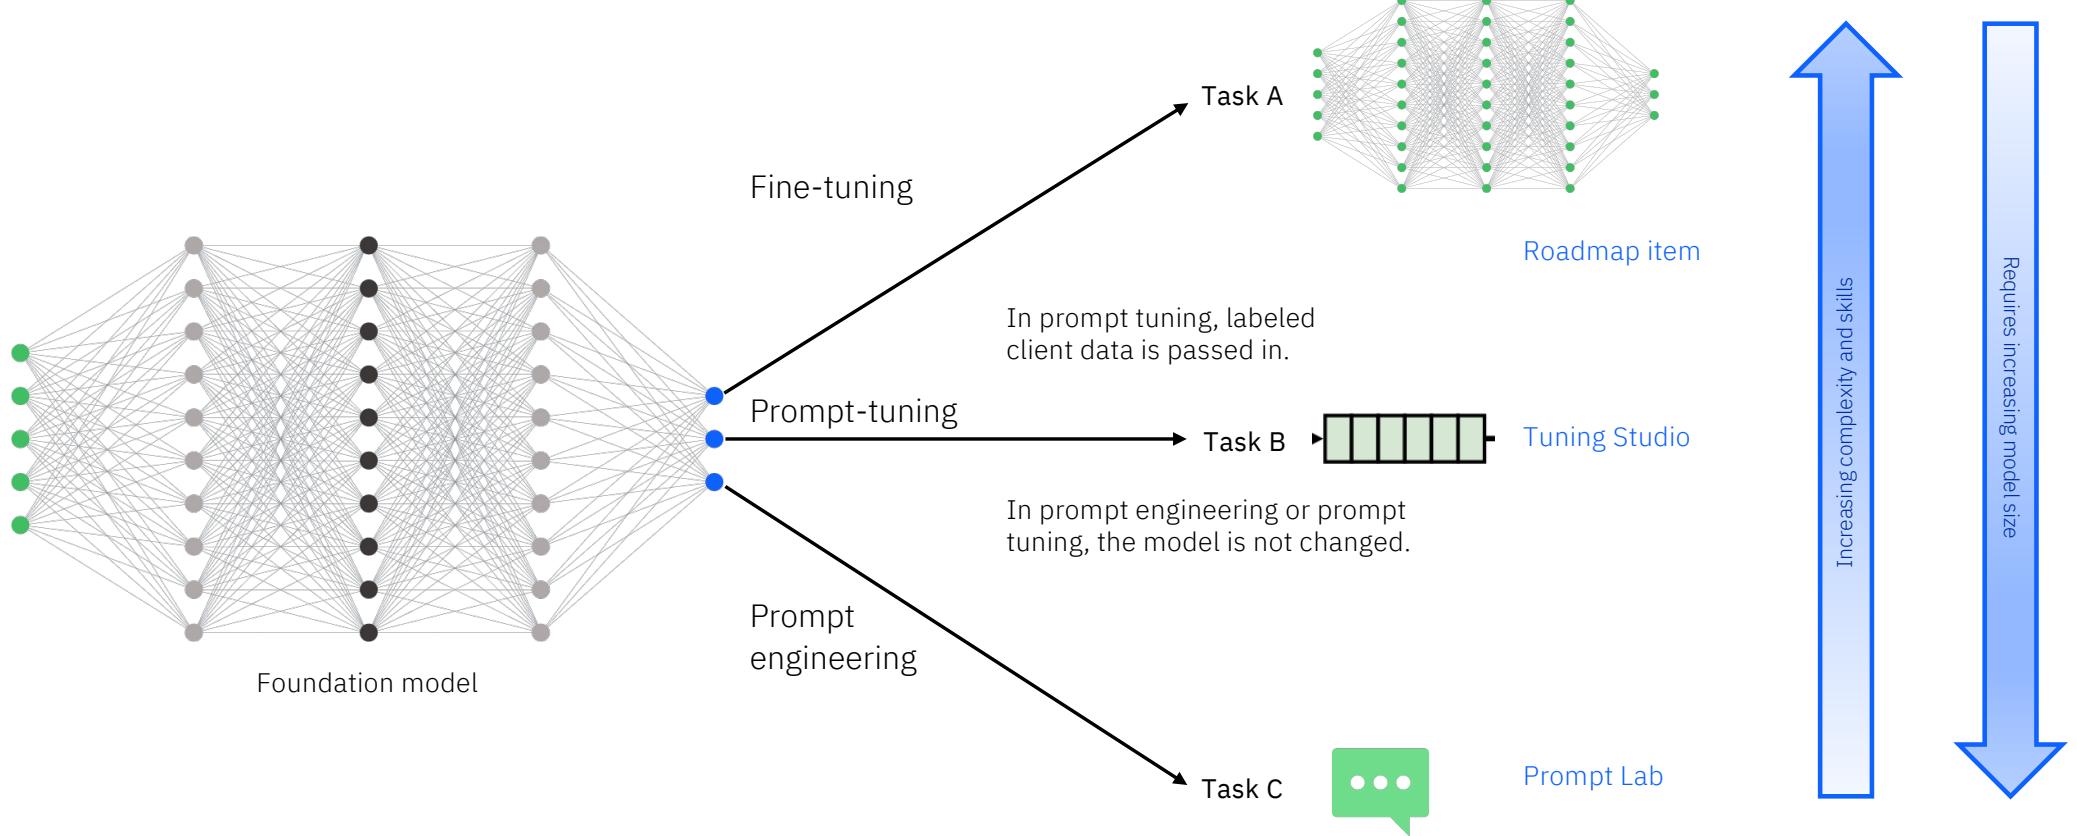
A type of foundation model trained with language-related data

ChatGPT is based on a large language model

# Foundation models: generalizable and adaptable



# Rapid adaptation to multiple tasks with small amounts of task-specific data



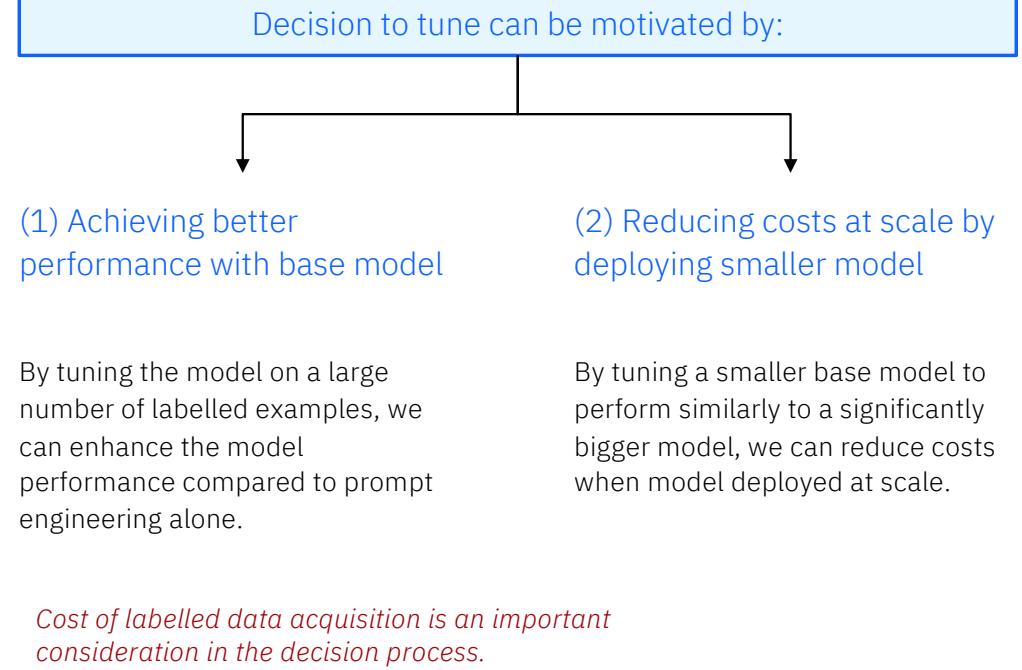
Fine-tuning requires labeled data and more resources to tune the model. When a model is fine-tuned, some of the weights are modified and clients get a private instance of the model.

# When to tune a model?

**Always start with prompt engineering** the largest LLM suited for your task.

This should provide some indication that the task is suited to be addressed by LLMs. It is also helpful to experiment with different labelled examples and understand which prompt formats work best on the target task.

7



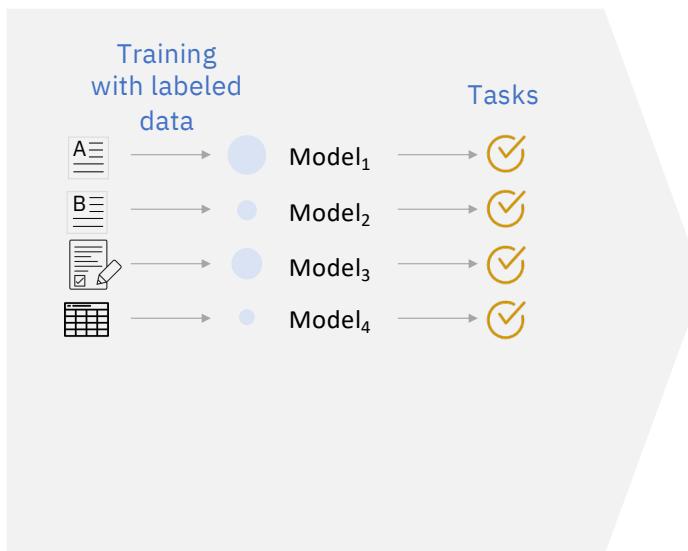
## Suggested workflow

	Step 1	Step 2	Step 3
Stage	Initial PoC	Pilot deployment	Deployment at scale
Goal	<i>Prove the use case with minimal effort</i>	<i>Reduce costs as permitted within PoC duration</i>	<i>Maximize ROI</i>
Recommendation	Use the <a href="#">largest model</a> and <a href="#">minimal labeled data</a>  Create labeled test dataset to measure model accuracy	Prompt engineer or prompt-tune a <a href="#">medium-size model</a> .  You may need to gather <a href="#">additional labeled data</a>	Consider using additional data gathered to fine-tune / prompt-tune <a href="#">a small model</a> .  Deploy the tuned model
Inference costs	\$\$\$	\$\$	\$

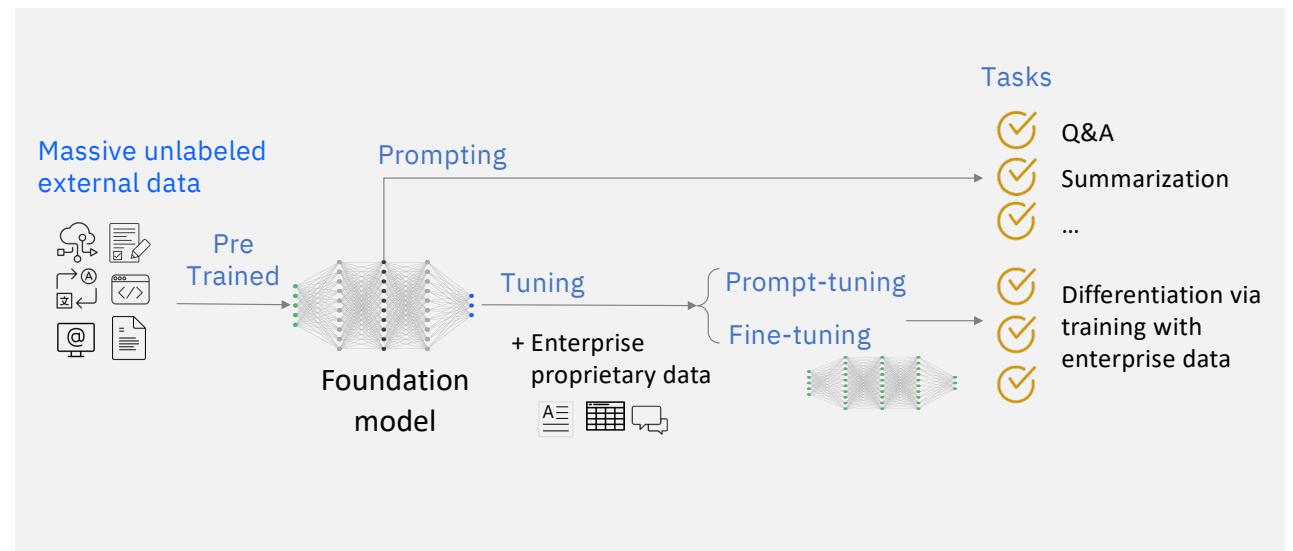
**Note:** fine-tuning a model requires creating a copy of the model specific to the user.  
The cost of hosting this model may impact the ROI analysis compared to prompt tuning.

# Foundational models enable a new paradigm of data-efficient AI development – generative AI

## Traditional AI models



## Foundation Models



- Individual siloed models
- Require task specific training
- Lots of human supervised training

- Rapid adaptation to multiple tasks with small amounts of task-specific data
- Pre-trained unsupervised learning

## IBM partnership with open-source models provider



**HUGGING FACE**

- IBM **watsonx.ai** clients have access to the latest and greatest open-source foundation models from Hugging Face.
- The IBM and Hugging Face partnership demonstrates a joint commitment to deliver an open ecosystem to clients, allowing them to find the best foundation models for their business needs.

# Most common generative AI tasks implemented today

Summarization	Classification	Generation
<p>Transform text with domain-specific content into personalized overviews that capture key points.</p> <p><i>Conversation summaries, insurance coverage, meeting transcripts, contract information</i></p>	<p>Read and classify written input with as few as zero examples.</p> <p><i>Sorting of customer complaints, threat and vulnerability classification, sentiment analysis, customer segmentation</i></p>	<p>Generate text content for a specific purpose.</p> <p><i>Marketing campaigns, job descriptions, blog posts and articles, email drafting support</i></p>

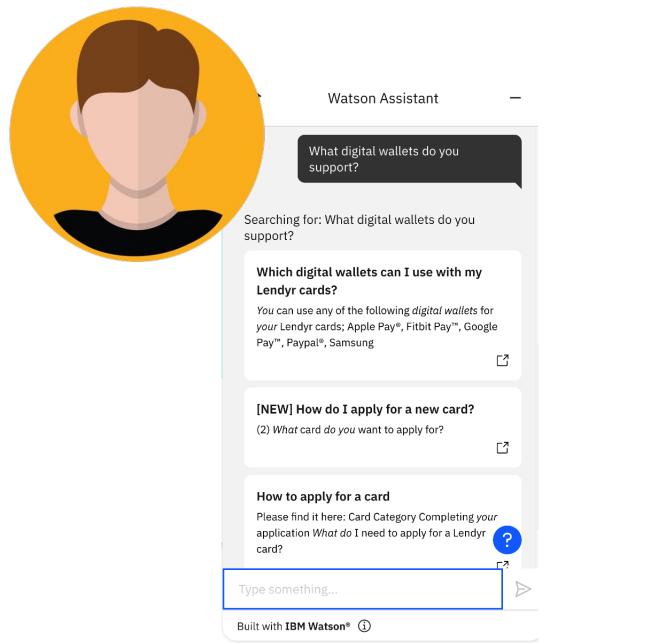
Extraction	Question-answering
<p>Analyze and extract essential information from unstructured text.</p> <p><i>Medical diagnosis support, user research findings</i></p>	<p>Create a question-answering feature grounded on specific content.</p> <p><i>Build a product specific Q&amp;A resource for customer service agents.</i></p>

# LangChain for LLMs

- LangChain is an open-source framework designed to simplify creating applications using LLMs.
- Models
- Prompt Templates
- Parsers
- Chains
- Question Answer

# Retrieval Augmented Search - Overview

## Conversational search – Q&A for documents



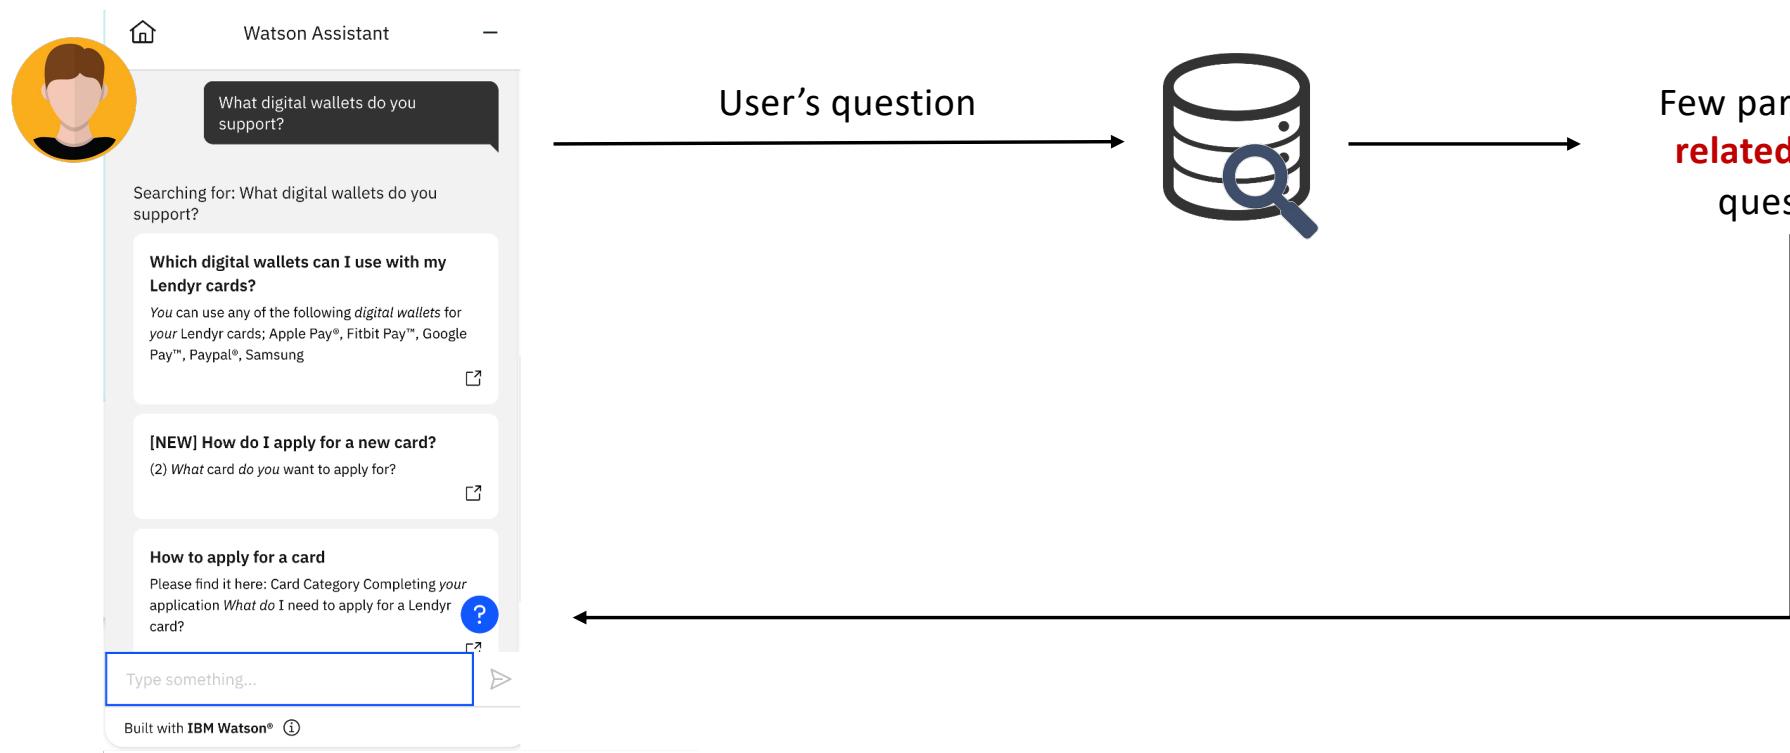
# Phase 1

Prepare the data



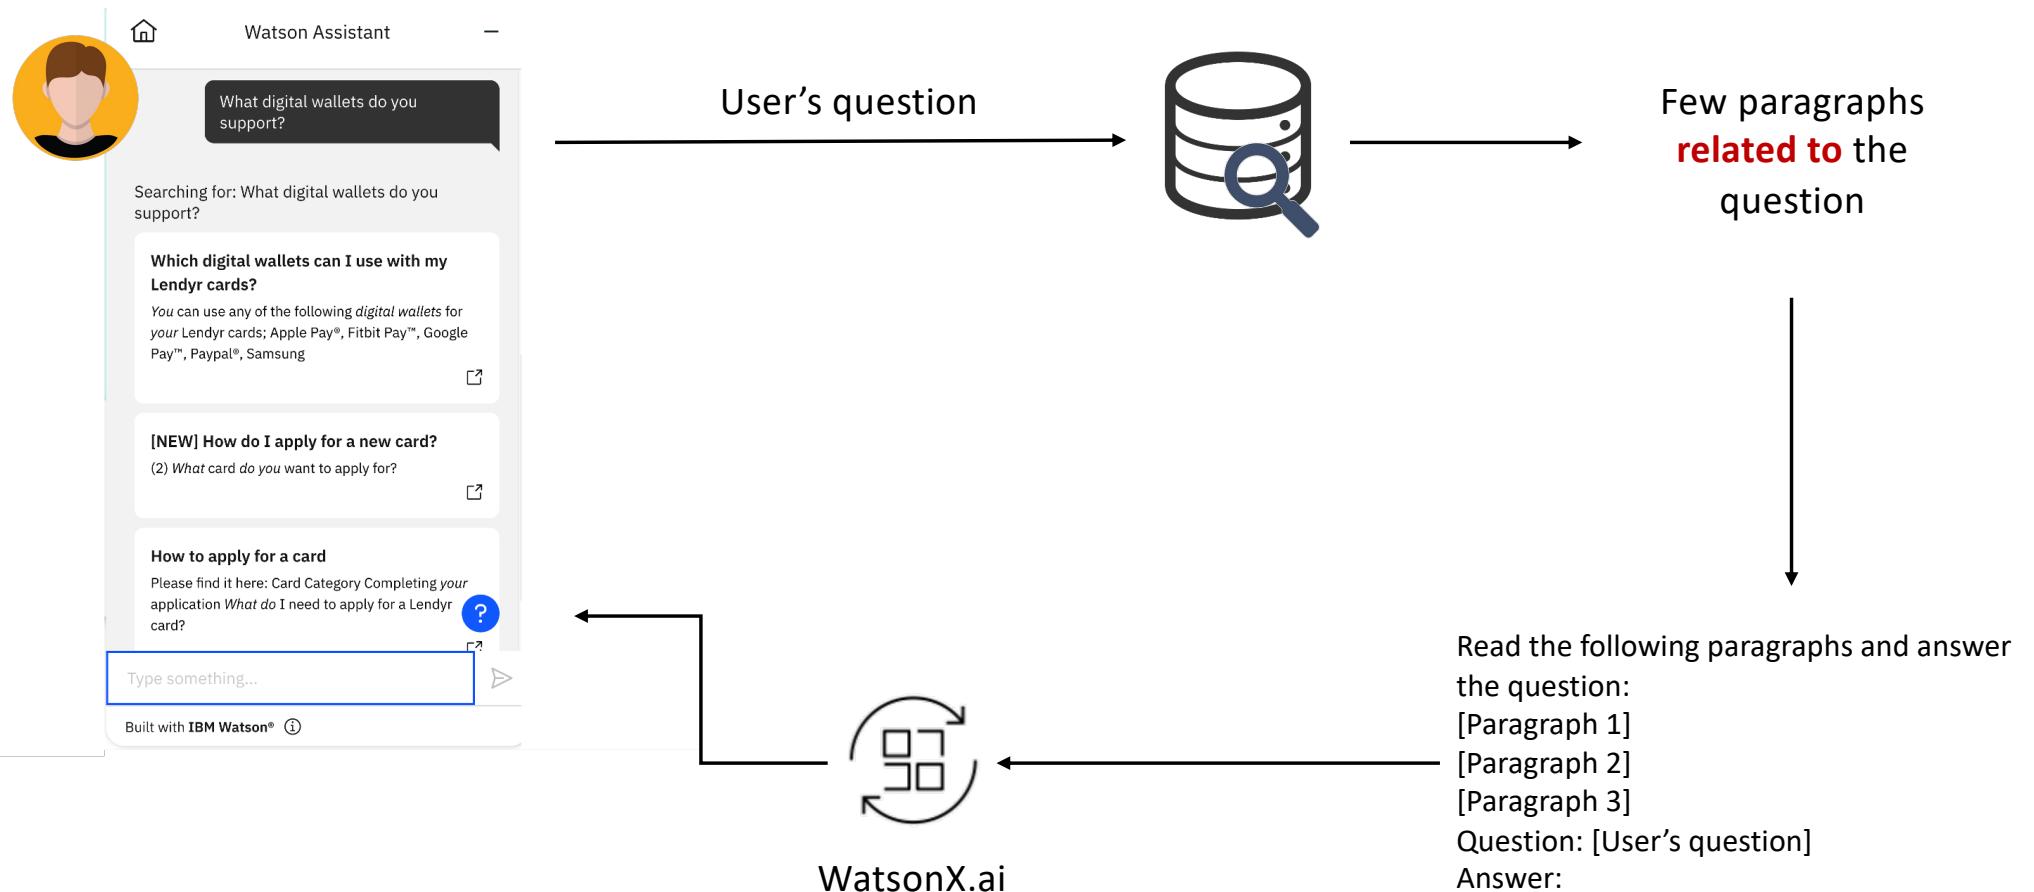
# Phase 2

## Query the data



# Phase 2 (new steps based on LLMs)

## Query the data

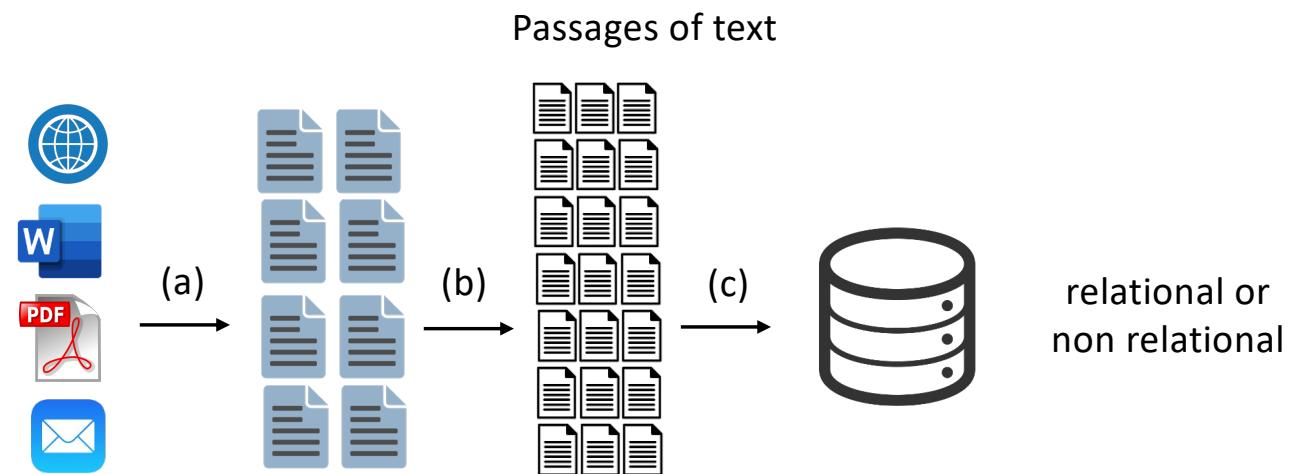


# Phase 1: The “traditional” way

## Phase 1

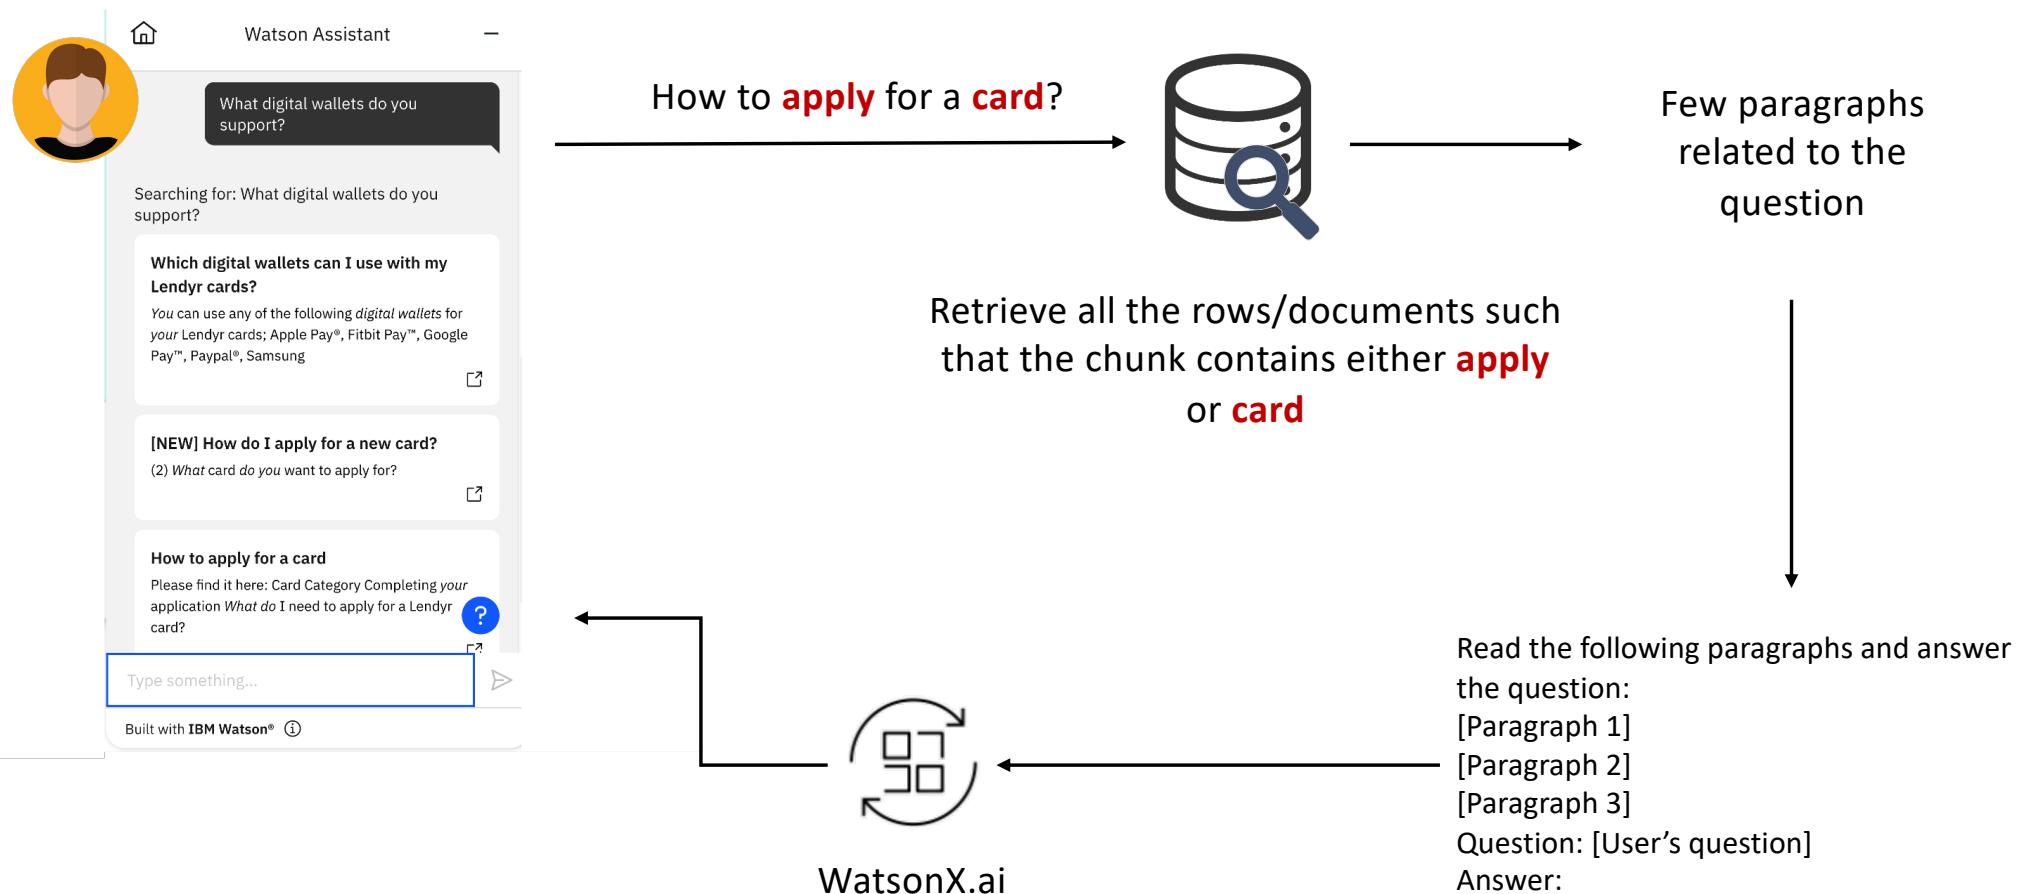
Ingest your data

- (a) Original files to documents
- (b) Documents to chunks
- (c) Chunks to database



# Phase 2: Syntactic search

## Query the data

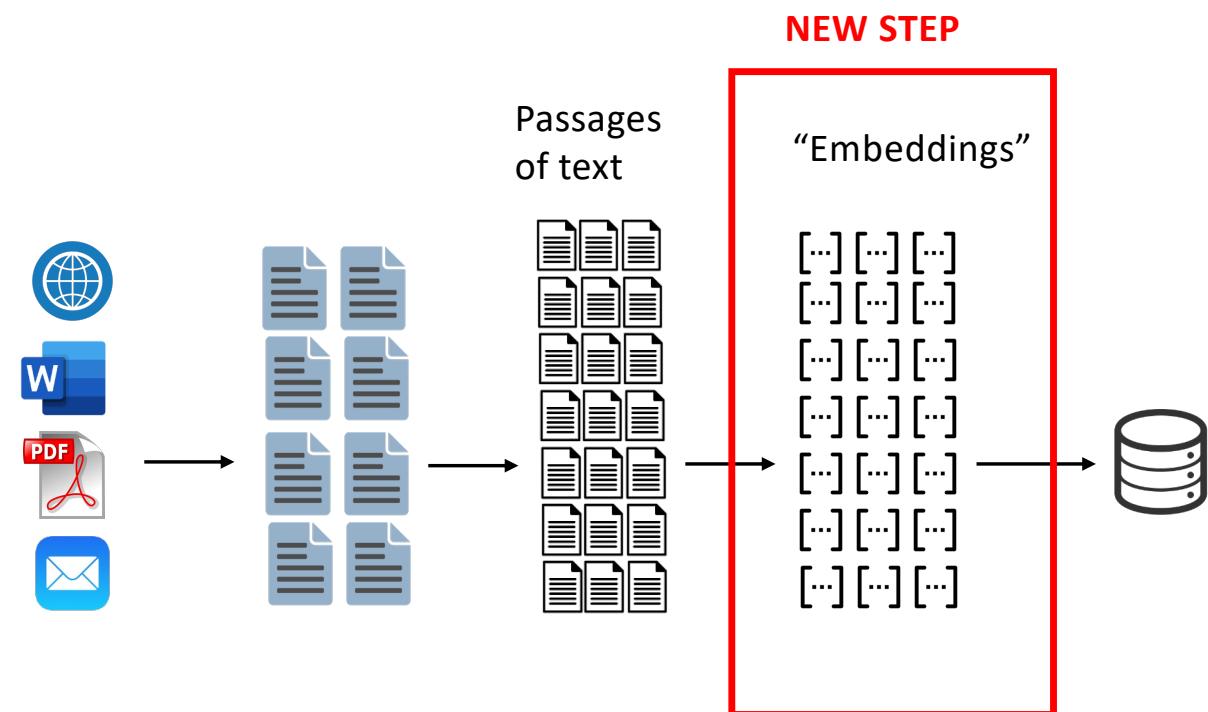


# Phase 1: The “embeddings” way

# Phase 1

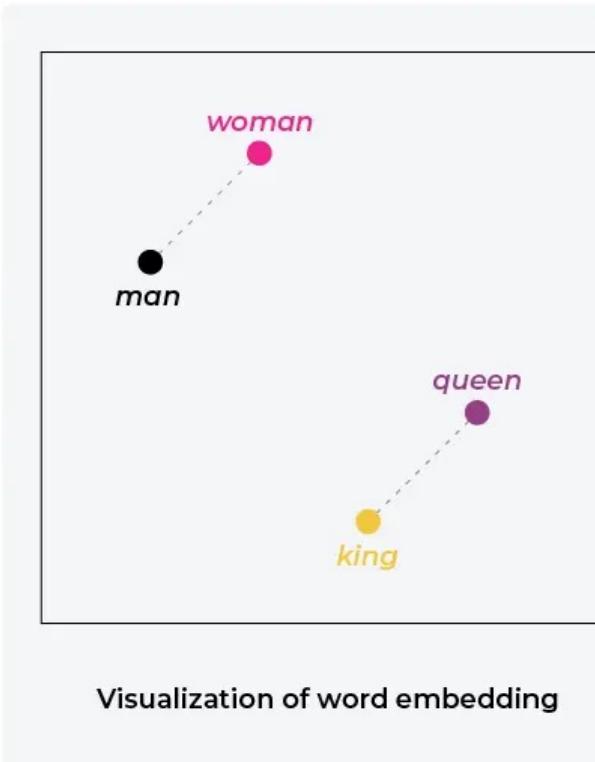
## Ingest your data

- (a) Original files to documents
  - (b) Documents to chunks
  - (c) Chunks to embeddings
  - (d) Embeddings to vector store



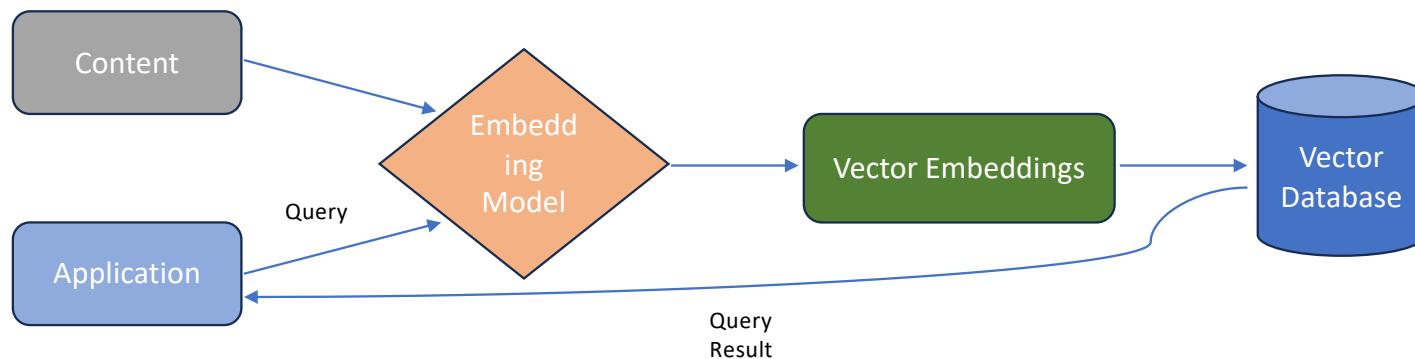
# Phase 1: The “embeddings” way

	word	living	being	feline	human	gender	royalty	verb	plural
<i>man</i>	→	0.6	-0.2	0.8	0.9	-0.1	-0.9	-0.7	
<i>woman</i>	→	0.7	0.3	0.8	-0.7	0.1	-0.5	-0.4	
<i>king</i>	→	0.5	-0.4	0.7	0.8	0.9	-0.7	-0.6	
<i>queen</i>	→	0.8	-0.1	0.8	-0.9	0.8	-0.5	-0.9	



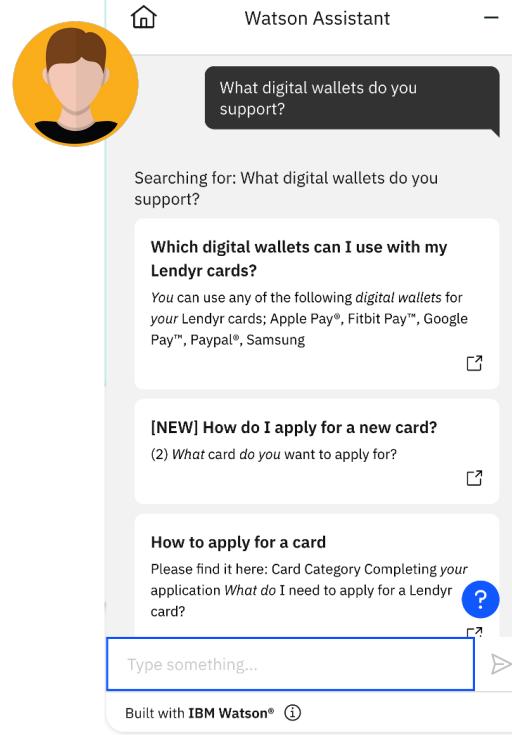
# Phase 1: The “embeddings” way – VectorDB

- A vector database is a special type of database that can store high-dimensional vectors which are mathematical representations of the features.
- This data is nothing, but a vector created through embeddings
- Use cases
  - Recommendation systems
  - Anomaly detections
  - NLP

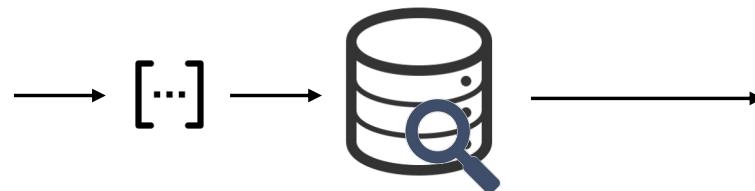


# Phase 2: Semantic search

## Query the data



How to apply  
for a card?



Retrieive all the embeddings that are  
**the closest to** the embedding of the  
user's question

Few paragraphs  
related to the  
question

Read the following paragraphs and answer  
the question:  
[Paragraph 1]  
[Paragraph 2]  
[Paragraph 3]  
Question: [User's question]  
Answer:

WatsonX.ai

# Syntactic vs. Semantic search

Why semantic is a “better” way of searching for information

The user expresses himself in his/her own way, whereas the documents usually use “specialized” terms.

Examples:



Paid leave of absence  
(IBM HR documents)

Corporate assets  
(Bank's code of ethics)

Revenue, profits, benefits  
(10K form)



Day off

Company's laptop

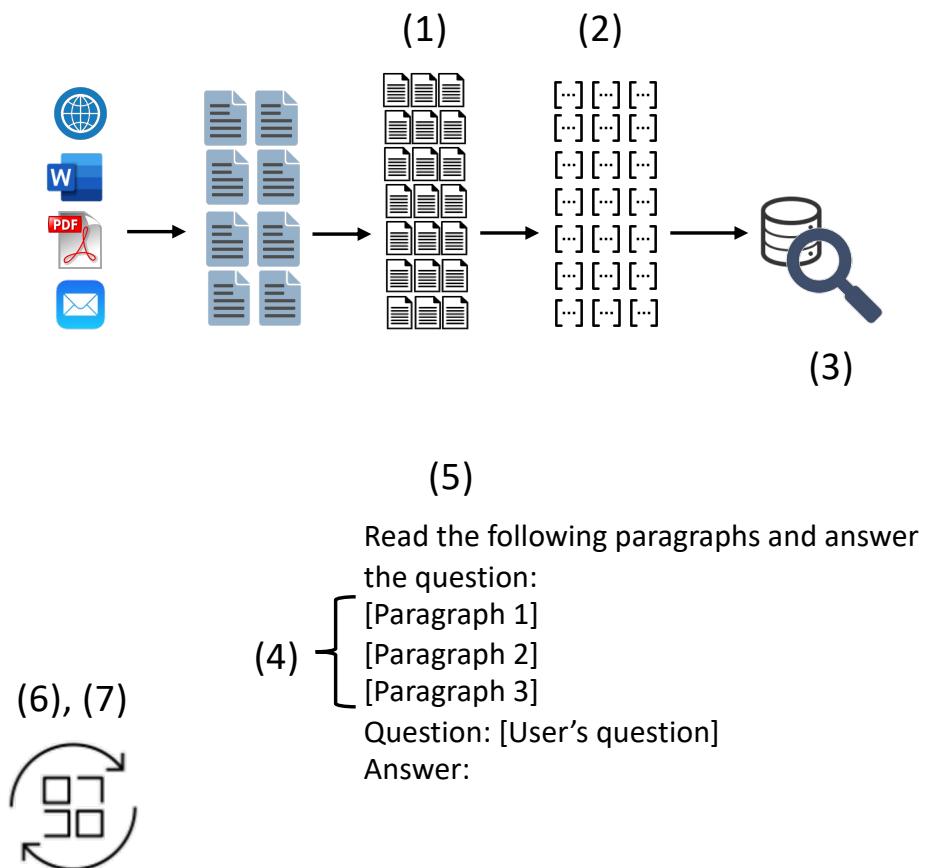
Money

# How to improve the accuracy?

Optimize the config at each and every step:

1. Length of the chunks of texts
2. Choice of embeddings library
3. Distance function between embeddings
4. Number of chunks retrieved from the database
5. Prompt
6. LLM parameters (temperature, topK, top, etc)
7. Choice of LLM
8. Etc.

**BUT there are more efficient ways**



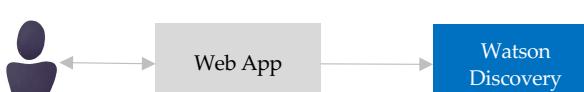
WatsonX.ai

# Watson / watsonx.ai Strawman Patterns

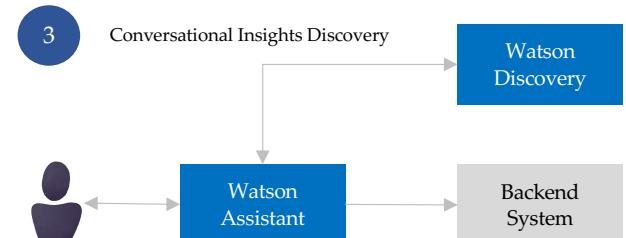
1 Classic Conversational AI



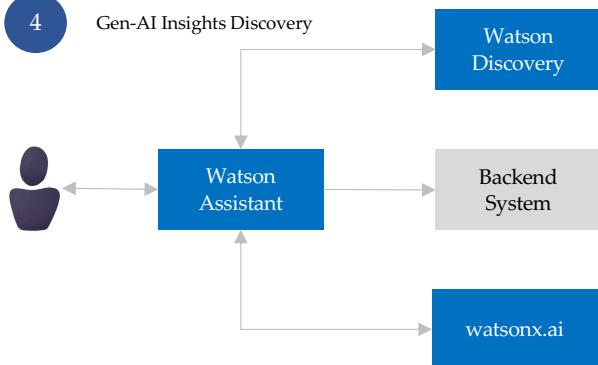
2 Classic Insights Discovery



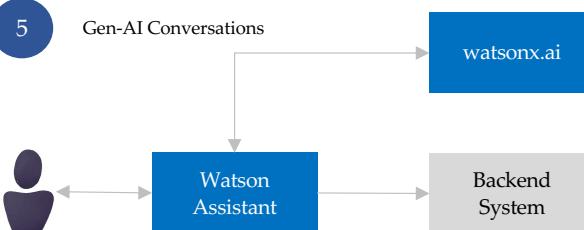
3 Conversational Insights Discovery



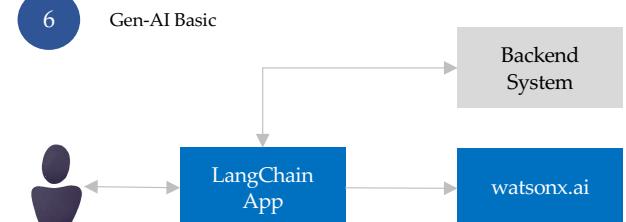
4 Gen-AI Insights Discovery



5 Gen-AI Conversations



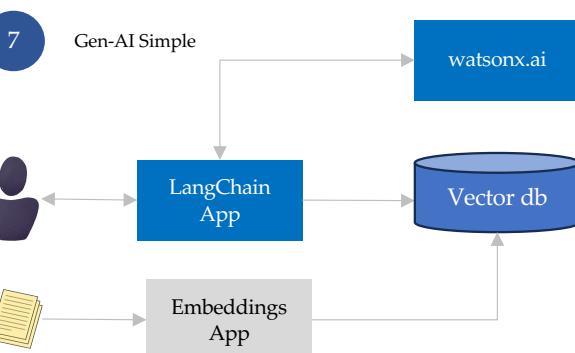
6 Gen-AI Basic



## Tasks

1. Summarization
2. Q & A
3. Extraction
4. Classification
5. Generation
6. Transformation
7. Listing
8. Comparison

7 Gen-AI Simple



## IBM Tech

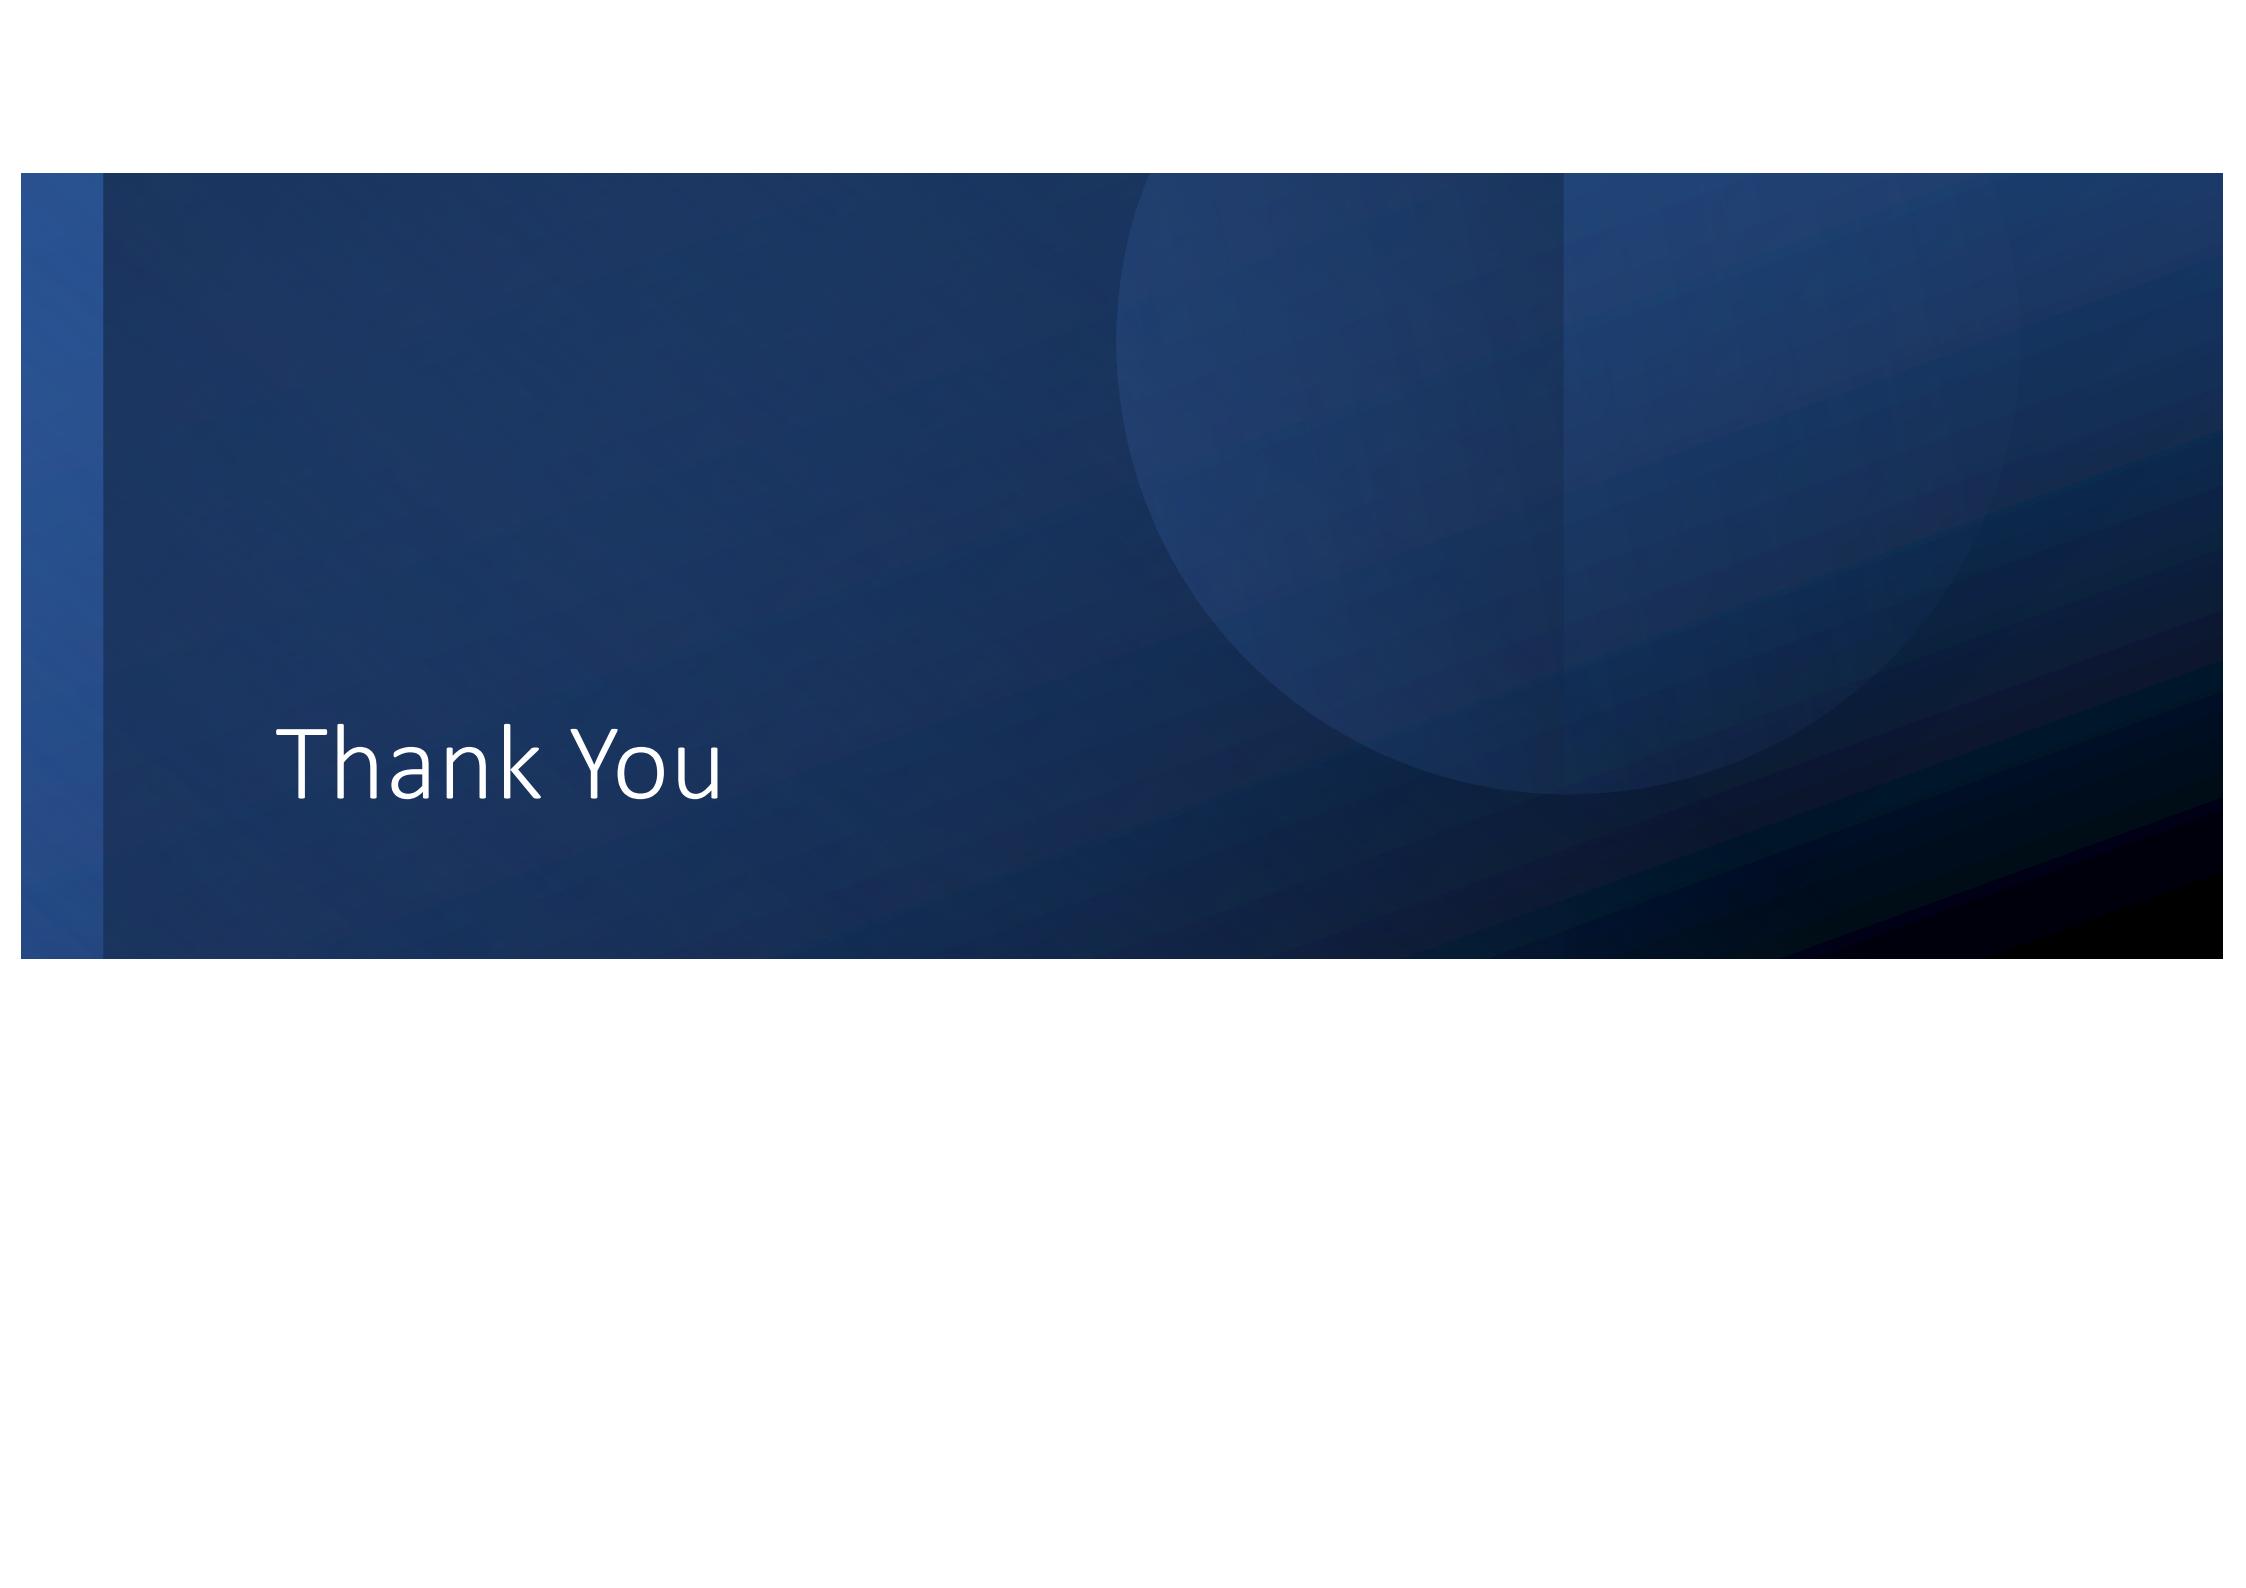
1. Foundation Models
2. Prompt Lab
3. Watson Assistant
4. Watson Discovery
5. Watson Speech
6. IBM Cloud

## Sample Technology Stack (non-IBM)

1. Python
2. LangChain
3. ChromaDB, Pinecone...
4. Flask, FastAPI...
5. MySQL

# Watsonx.ai using Python SDK

## Demo



Thank You